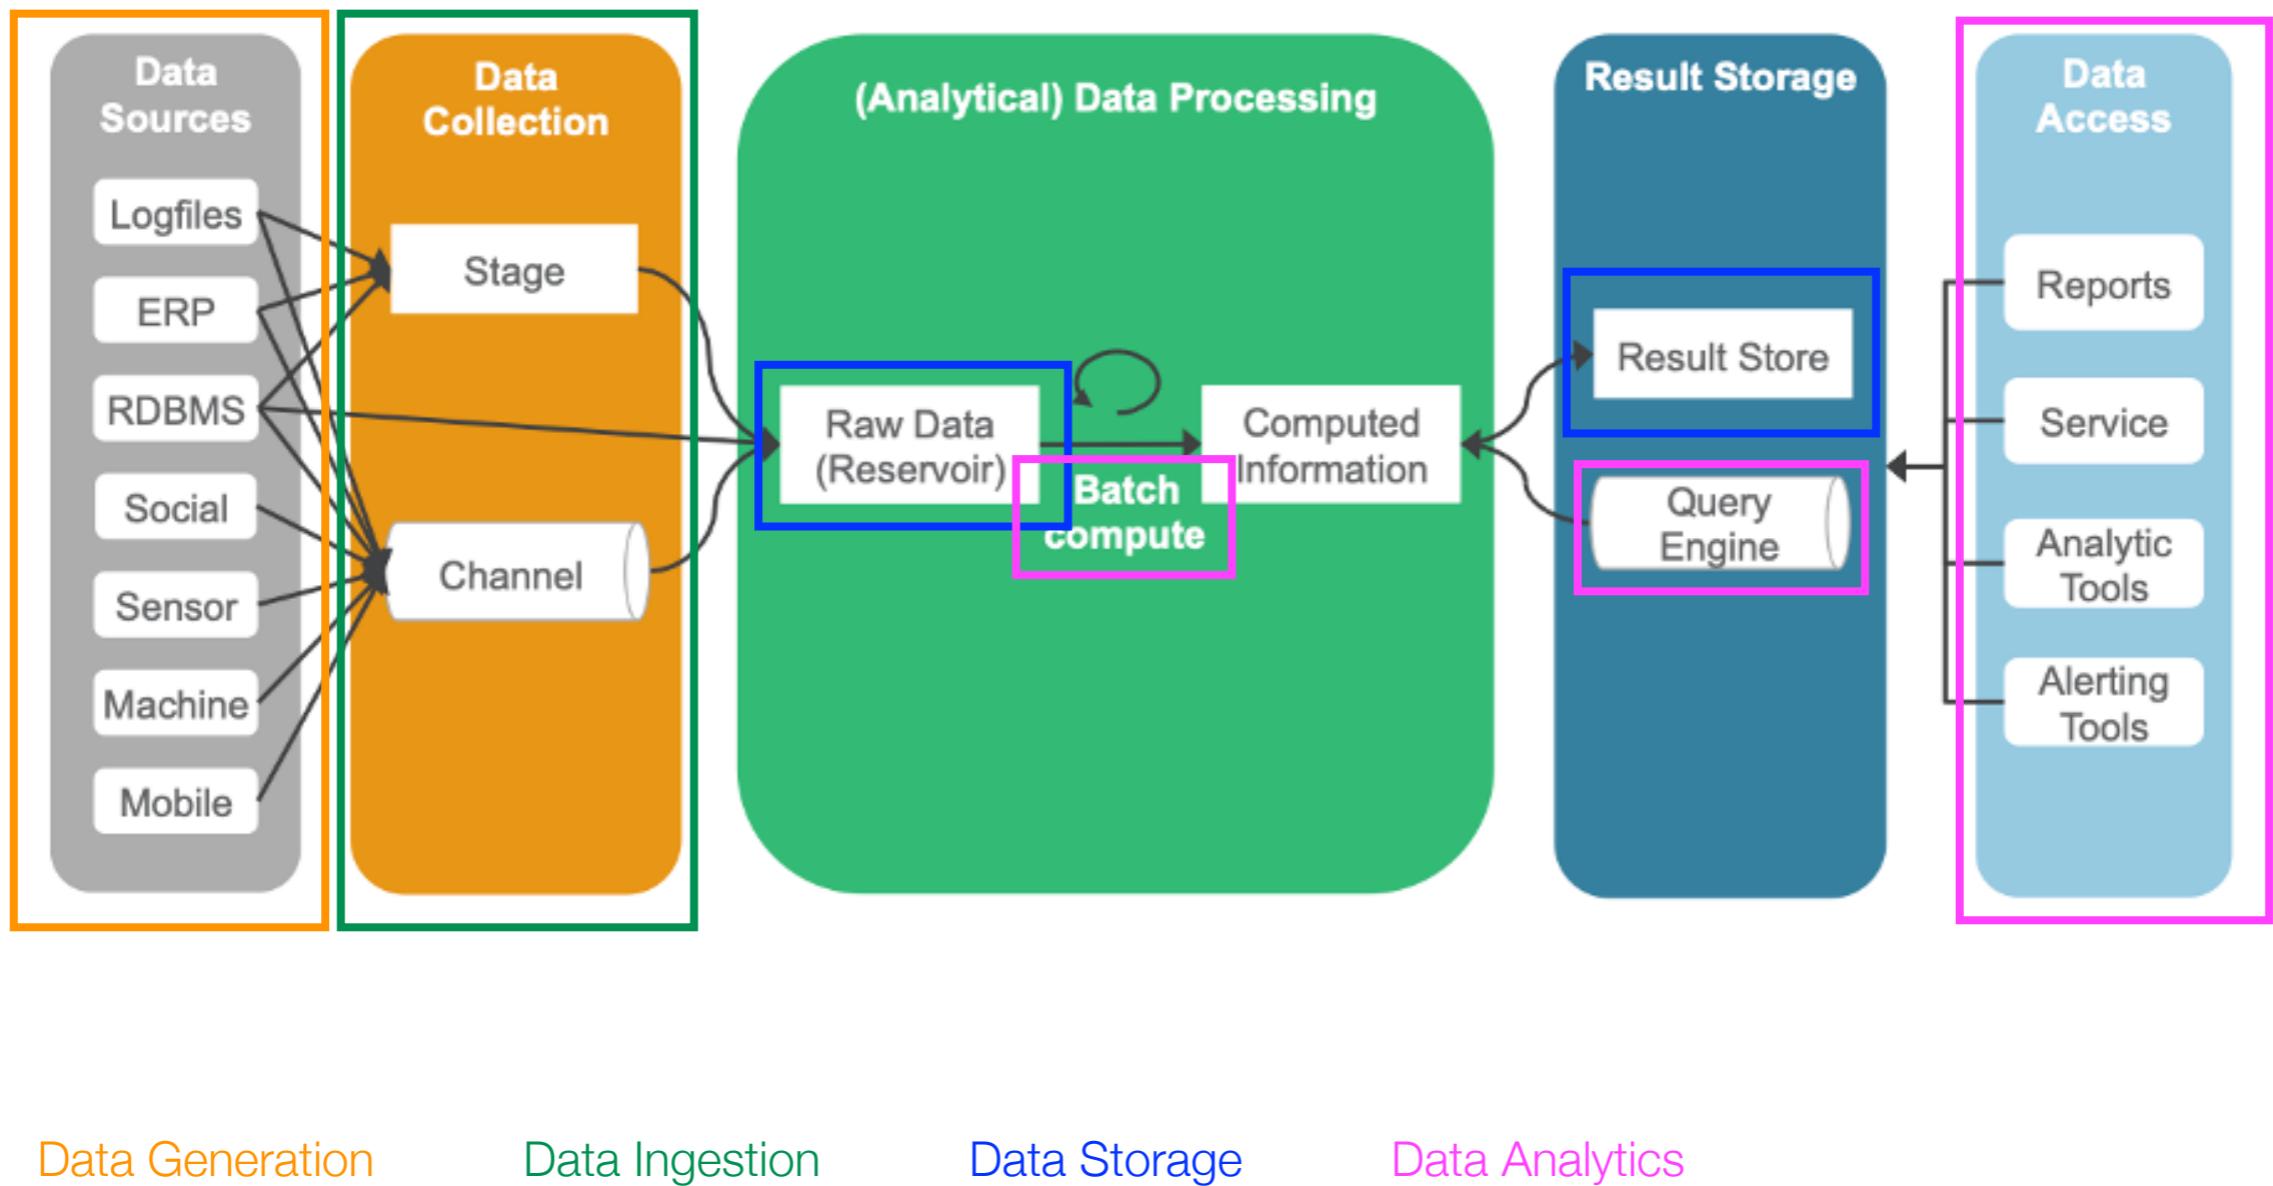


2110531 - Data Science and Data Engineering Tools

Data Storages

Asst.Prof. Natawut Nupairoj, Ph.D.
Department of Computer Engineering
Chulalongkorn University
natawut.n@chula.ac.th

Simple Big Data Analytic Architecture





Traditional Database

- Based on "relational model"
 - Data is split and stored into tables
 - Tables can be processed together using set-like operations
 - Data model is usually normalized to remove duplication
- Very suitable for OLTP or transaction systems
 - Provide lots of complicated SQL operations
 - Lots of inserts and updates

Problems of Data Science Storage

- There are several needs for data analytics purposes e.g. traditional data store, caching, feature store
- Data is historical data and its volume can be huge
- Scalability is extremely important and “Relational + Consistency” can limit scalability
- SQL command can be very complex and time-consuming
 - It requires the synchronization of data accessing between multiple tables
 - It will be poor when using on more than a few servers in the same cluster

Relational Database: Normalized Data Model




How can we split these tables to lots of machines?


- Dishes of the same category in the same machine
- How about dishes and ingredients?

Or we can replicate data – lead to data consistency problems


Replication and Data Consistency Problem



Solution to Data Consistency - Active Replication



Solution to Data Consistency - Passive Replication



Data Consistency Models

- Strong consistency
 - After update completes, any access will return the updated value
 - High costs for large scale system
- Weak consistency
 - Certain conditions must be met before the consistency is guaranteed
 - Time between update completion and guaranteed consistency is called **inconsistency window**

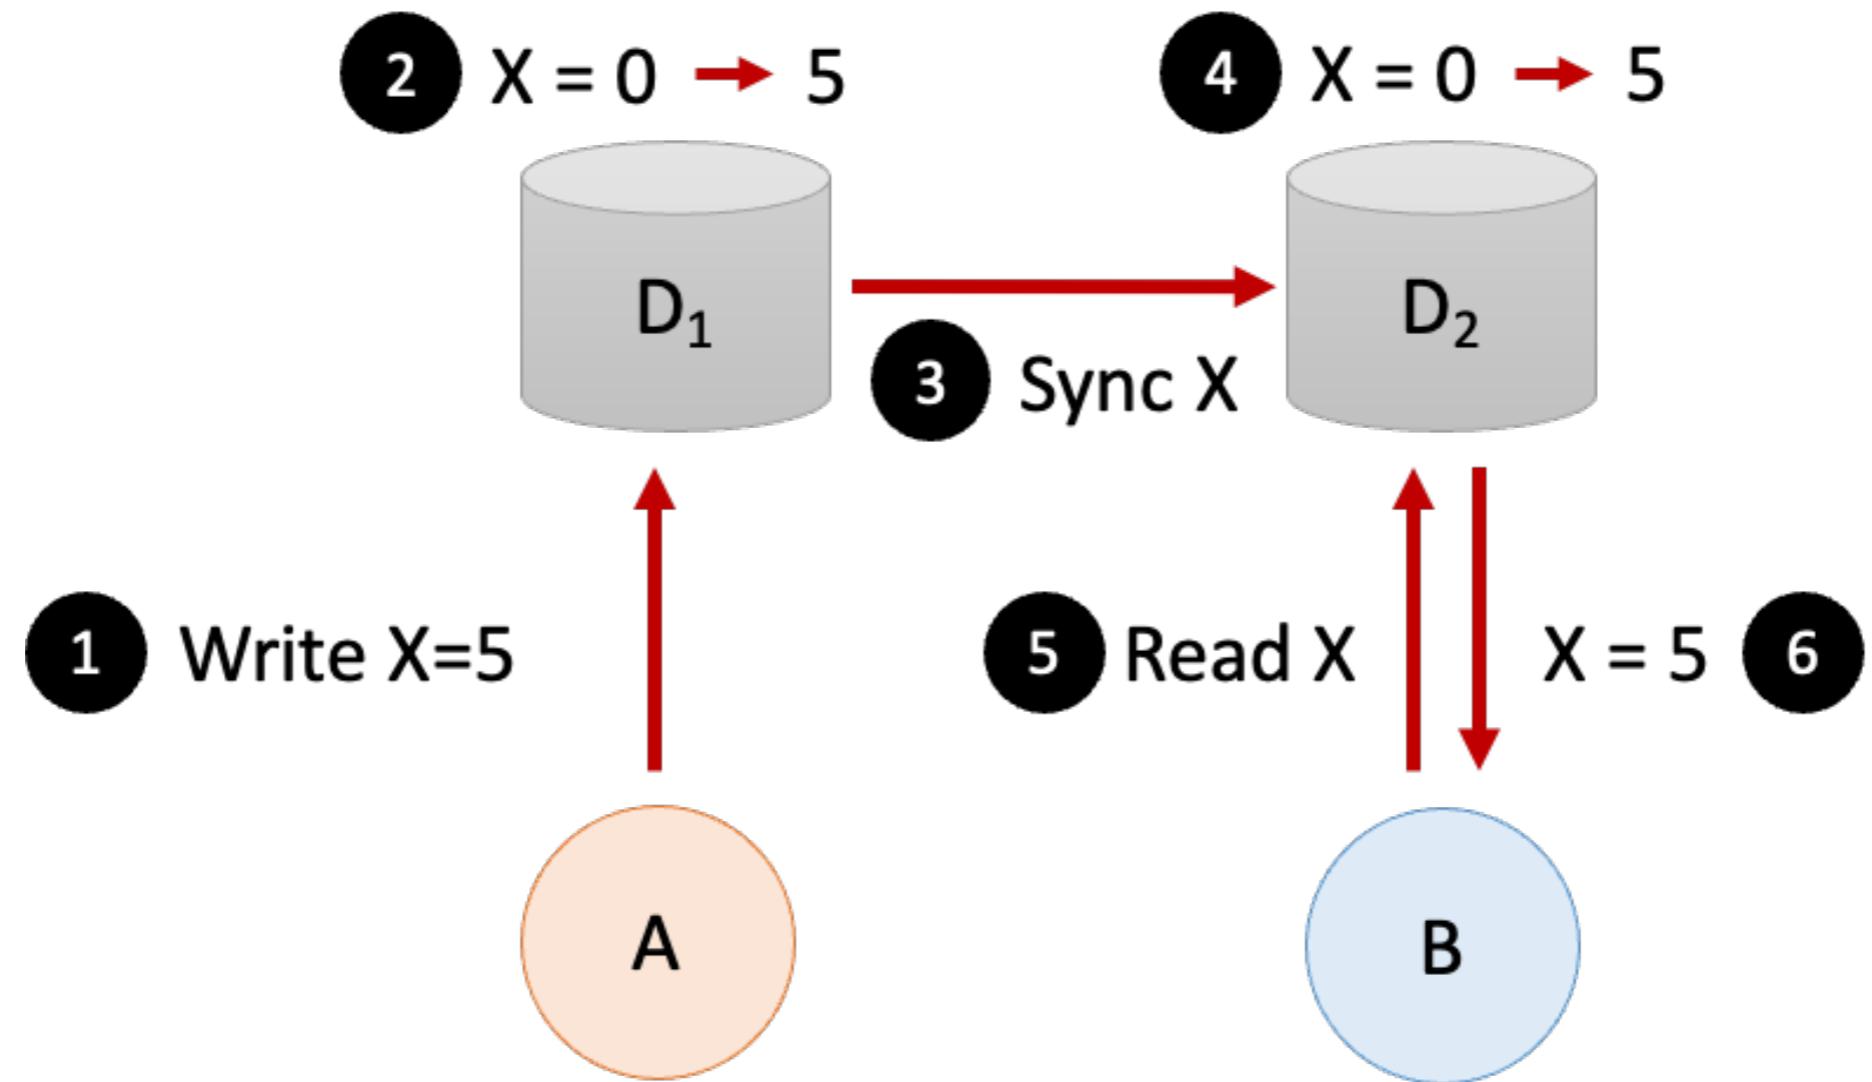
CAP Theorem (Brewer's Theorem)

- By Eric Brewer (University of California, Berkeley)
- It is impossible for a distributed computer system to simultaneously provide all three of the following guarantees: Consistency, Availability, Partition tolerance
- Scenario
 - Distributed system (clients and servers)
 - Multiple servers working together
 - Multiple clients may read or write on the same data at the same time


CAP Guarantees only 2 out of 3

- Consistency
 - Every read receives the most recent write or an error
- Availability
 - Every request receives a (non-error) response – without guarantee that it contains the most recent write
 - Response from any server is good
- Partition tolerance
 - The system continues to operate despite arbitrary message loss or failure of part of the system
 - Lots of servers require long synchronization time causing servers to not be able to communicate among one another within reasonable time

CA Data Store



Given P, Choosing between C and A (at step 5)




The Landscape of NoSQL

- Alternatives to SQL database - non-relational, distributed, and horizontally scalable
- Data is shared and distributed across multiple servers
- Typically use weak consistency model (but not always)
- Examples
 - Document: MongoDB, DynamoDB, CosmosDB, Couchbase, Firebase
 - Column: Cassandra, HBase, CosmosDB, Accumulo
 - Key-value: Redis, DynamoDB, CosmosDB, MemcacheDB
 - Graph: Neo4J, CosmosDB, ArrangoDB, OrientDB
 - Search Engine: Elasticsearch, Splunk, Solr

How NoSQL can “Scale”

- Principle ideas
 - Split data into chunks or shards
 - Distribute data across multiple servers
 - Must require minimum synchronization
- Have to give up some traditional features
 - No complex relational model
 - Relax consistency
 - Duplicated information (not space optimized)
 - Fast to insert new record, but not so fast to update the existing one



Source: <https://dev.to/katkelly/cap-theorem-why-you-can-t-have-it-all-ga1>

Redis

Key-Value Store

Data Storage




Key-Value Stores

- A data store designed for storing, retrieving, and managing associative arrays (aka. dictionary or hash table)
- Main concept is to store data as a collection of key-value pairs in which a key serves as a unique identifier
- Simple and fast, often use in-memory architecture, ability to scale-out
- Use cases: session store, shopping cart, SQL/API caching, task queue, etc.

Redis (Remote Dictionary Server)


- In-memory data structure store with clustering, transactional, time-to-live limiting, and auto-failover capabilities
- Support wide-range of data structure with lots of related operations for each structure
- Being used for database cache, message broker, streaming engine, feature store engine, etc.
- Provide CLI and support many programming languages
- Many useful modules are available to extend the functionality of Redis core e.g. RedisJSON, RedisSearch, RedisTimeSeries, Redis OM, etc.

How is redis traditionally used



Redis Persistence

- Redis supports many level of persistence: no persistence, RDB (point-in-time snapshot), AOF (log every write), RDB+AOF



Running Redis

- The simplest way to run a redis instance is to use docker

```
docker pull redis
docker run -d --rm --name redis -p 6379:6379 redis
```
- This will start redis in your docker at port 6379 and map the port to your localhost
- You can also use docker-compose.yml and other example files in redis folder in datasci_architecture repo

Working with Redis


- Redis CLI
 - Standard client program to connect to any redis server
 - Come with any redis installation (see: <https://redis.io/docs/getting-started/>)

redis-cli

redis-cli -h lab.aimet.tech

- You can type in redis command in the CLI input

Redis Data Types



String

- Similar to Python or Java Strings, maximum length of 512MB

```
SET "user" "Natawut Nupairoj"
```

```
GET "user"
```

```
DEL "user"
```

- Use cases
 - Server-side object cache e.g. HTML fragments, shopping cart, user profile
 - Queues
 - Activity tracking

Other String Commands

APPEND

INCR

SET

DECR

INCRBY

SETEX

DECRBY

INCRBYFLOAT

SETNX

GET

LCS

SETRANGE

GETDEL

MGET

STRLEN

GETEX

MSET

SUBSTR

GETRANGE

MSETNX

GETSET

PSETEX

Useful Commands

- Any item in Redis can be made to expire after or at a certain time

```
EXPIRE user 60. # in seconds
```

```
TTL user
```

- You can scan all index with scan command

```
SCAN 0
```

- You can delete item or test its existence

```
DEL mykey
```

```
EXISTS mykey
```

List

- List of strings, sorted by insertion order

- Can be used as list, queue, stack

```
LPUSH mylist abc # mylist contains "abc"
```

```
LPUSH mylist xyz # mylist contains "xyz", "abc"
```

```
RPUSH mylist 123 # mylist contains "xyz", "abc", "123"
```

- Use cases

- Queue

- Timelines

List Commands

BLMOVE

LMOVE

LSET

BLMPOP

LMPOP

LTRIM

BLPOP

LPOP

RPOP

BRPOP

LPOS

RPOPLPUSH

BRPOPLPUSH

LPUSH

RPUSH

LINDEX

LPUSHX

RPUSHNX

LINSERT

LRANGE

LLEN

LREM

Set

- Powerful data types for unordered non-duplicated keys
- Support many set operations e.g. intersection, union, etc.

```
SADD user_set natawut
SCARD user_set
SMEMBERS user_set
```
- Use cases
 - Set of user profiles
 - Set of inappropriate words for inappropriate content filtering

Set Commands

SADD

SISMEMBER

SSCAN

SCARD

SMEMBERS

SUNION

SDIFF

SMISMEMBER

SUNIONSTORE

SDIFFSTORE

SMOVE

SINTER

SPOP

SINTERCARD

SRANDMEMBER

SINTERSTORE

SREM

Sorted Set

- Set of sorted items based on the score associated to each member

```
ZADD my_sortedset 5 data1
```

```
ZADD my_sortedset 1 data2 10 data3
```

```
ZRANGEBYSCORE my_sortedset 5. +inf WITHSCORES
```

- Use
 - Leader scoreboard
 - Priority queue

Sorted Set Commands

BZMPOP

ZDIFFSTORE

ZMSCORE

BZPOPMAX

ZINCRBY

ZPOPMAX

BZPOPMIN

ZINTER

ZPOPMIN

ZADD

ZINTERCARD

ZRANDMEMBER

ZCARD

ZINTERSTORE

ZRANGE

ZCOUNT

ZLEXCOUNT

ZRANGEBYLEX

ZDIFF

ZMPOP

ZRANGEBYSCORE

Hash

- A container of unique fields and their values

```
HMSET profile:12345 user nnp id 12345 name "Natawut Nupairoj"  
balance 10
```

```
HGETALL profile:12345
```

```
HINCRBYprofile:12345 balance 5
```

- Use
 - User profile information
 - Post. Information

Hash Commands

HDEL

HLEN

HSTRLEN

HEXISTS

HMGET

HVALS

HGET

HMSET

HGETALL

HRANDFIELD

HINCRBY

HSCAN

HINCRBYFLOAT

HSET


HKEYS

HSETNX

Example: Recipes in Redis

- Create data model for recipe
- Data items
 - Main data
 - Dish - hash - key = dish:dish_id, value = dish_id, name, category, cuisine
 - Dish Ingredients - set - key = d_ing:dish_id, value = string(name; size; unit)
 - For query
 - Cuisine - set - key = cuisine:cuisine_value, value = string(dish_id)
- Relationship
 - 1 dish can have multiple ingredients
 - Cuisine, as a helper data item, points back to dishes belong in the cuisine
- This is only one example of model design

Example: Recipes in Redis



```
HMSET dish:1 id 101 name "Pad Thai" cuisine "thai" category "noodle"
```

```
SADD d_ing:101 "noodle;125;g"
```


```
SADD d_ing:101 "shrimp;150;g"
```

```
SADD d_ing:101 "brown sugar;3;tbsp"
```

```
SADD cuisine:thai "101"
```

Redis Architecture

- Redis supports flexible architecture
 - Single instance
 - High Availability
 - Sentinel
 - Cluster



Source: <https://architecturenotes.co/redis/>


Redis High Availability

- When setup for HA, Redis consists of a master and one or more replicas and follows the following rules
 - Master performs all commands (read/write/etc.), replica performs only read commands
 - When a master and a replica instances are well-connected, the master keeps the replica updated by sending a stream of commands to the replica to replicate the effects on the dataset
 - When the link between the master and the replica breaks, for network issues or because a timeout is sensed in the master or the replica, the replica reconnects and attempts to proceed with a partial resynchronization
 - When a partial resynchronization is not possible, the replica will ask for a full resynchronization
 - Redis uses asynchronous replication, with asynchronous replica-to-master acknowledges of the amount of data processed

Redis High Availability

- To ensure data synchronization, all servers keep data version (replication id, offset); equal data versions mean both servers hold the same data
 - When a server becomes a master, it generates a new replication ID; replica server inherits replication ID from its master once connected
 - When master updates its data, it increments offset by one
- When replicas connect to masters, they use the PSYNC command to send their old master replication ID and the offsets they processed so far
 - If the differences of offsets are not too far behind, partial synchronization is performed
 - If too far or mismatched replication ID, a full resynchronization happens: in this case the replica will get a full copy of the dataset, from scratch

Redis Sentinel




- A cluster of sentinel processes working together to coordinate states of Redis master (main) and replicas (secondaries)
- Focus on HA aspects with minimum number of nodes
- Support single master with multiple replicas with monitoring, notification, auto-failover, configuration providing
- Need at least three Sentinel instances for a robust deployment (quorum or agreement protocol)
- Can be deployed in Redis and client nodes

Source: <https://architecturenotes.co/redis/>

Redis Cluster

- Support algorithmic sharding for horizontal scaling
 - Incoming data is hashed into a hash slot using deterministic hash function and mod with number of nodes
 - Data with a given key always maps to the same shard
 - Perform resharding when adding new nodes
- There are 16k hashslot; all data in the same hashslot always moves to the new node when resharding
- Nodes communicate with one another using gossiping protocol to determine the entire cluster's health
- Redis Cluster does not guarantee strong consistency



Source: <https://architecturenotes.co/redis/>

Redis in Python

- Redis-Py
 - Standard python package for redis client

pip install redis

- Example


```
import redis

r = redis.Redis(host='hostname', port=port)
```

Example: Simple Social Network

- Create data model for simple social network
- Data items
 - Users - id, name, can follow others, can be followed
 - Posts - id, content
- Relationship
 - 1 user can have many posts, each post can associate to only one user
 - User can follow one another

Users and Posts




```
SET user:101:name pinot           # user id 101 with name "pinot"  
SET username:pinot 101            # refer back to user-id
```

```
HMSET post:1 user 101 content "this is the first post"
```

```
RPUSH user:101:post 1  
RPUSH user:101:post 8
```

Users and Followers



```
SADD user:101:follows 104
```

```
SADD user:101:follows 105
```

```
SADD user:104:followed_by 101
```

```
SADD user:105:followed_by 101
```

Simple Redis-Py Example

This notebook contains simple redis python commands.

For your local redis only

In [1]: `import redis`

Connect to local server -- no hostname or ip is needed

In [2]: `rd = redis.Redis(charset="utf-8", decode_responses=True)`

In [3]: `rd.set('user:101:name', 'pinot')`

Out[3]: True

In [4]: `rd.get('user:101:name')`

Out[4]: 'pinot'

In [5]: `rd.hset('post:1', 'user', 101)
rd.hset('post:1', 'content', 'this is the first post')`

Out[5]: 1

In [6]: `rd.hgetall('post:1')`

Out[6]: {'user': '101', 'content': 'this is the first post'}

In [7]: `rd.rpush('user:101:post', 1)
rd.rpush('user:101:post', 8)`

Out[7]: 2

In [8]: `rd.llen('user:101:post')`

Out[8]: 2

```
In [9]: rd.lrange('user:101:post', 0, -1)
```

```
Out[9]: ['1', '8']
```

```
In [10]: rd.sadd('user:101:follows', 104)
rd.sadd('user:101:follows', 105)
```

```
Out[10]: 1
```

```
In [11]: rd.scard('user:101:follows')
```

```
Out[11]: 2
```

```
In [12]: rd.smembers('user:101:follows')
```

```
Out[12]: {'104', '105'}
```

```
In [13]: cursor = 0
cursor, keys = rd.scan(cursor=cursor, match='user:*')
while cursor > 0:
    for key in keys:
        print('found: ', key)
    cursor, keys = rd.scan(cursor=cursor, match='username:*')

for key in keys:
    print('found: ', key)
```

```
found: user:101:follows
found: user:101:name
found: user:101:post
```

Assignment: Simple Social Network

- Use the redis assignment notebook in class git repo and provide the following answers:
 - Answer questions in the question set
 - Save your notebook solution in the PDF format similar to the example in class git repo (used File -> Print Preview and then export as PDF or print to PDF with your browser)

My Work

Attach file(s) Answer a question set

Complete the question set

1 What is the username of user id "600"?
Pick a choice:

giddyCheetah7

humorousDinosaur4

cautiousCrackers9

yearningPaella1

grizzledOryx9

grudgingMussel2

panickyBaboon8

2 What is the id of username "excitedPie4" ?
Pick a choice:

125

References

- ScaleGrid, “Top Redis Use Cases by. Core Data Structure Types”, <https://scalegrid.io/blog/top-redis-use-cases-by-core-data-structure-types/>
- Jerry An, “The most important Redis data structures you must understand”, <https://medium.com/analytics-vidhya/the-most-important-redis-data-structures-you-must-understand-2e95b5cf2bce>
- Brad Solomon, “How to use Redis with python”, <https://realpython.com/python-redis/>
- M. Yusuf, “Redis Explained”, <https://architecturenotes.co/redis/>